



Space-based faults




Daniel Krolopp | Aaron Neustetter



Motivation

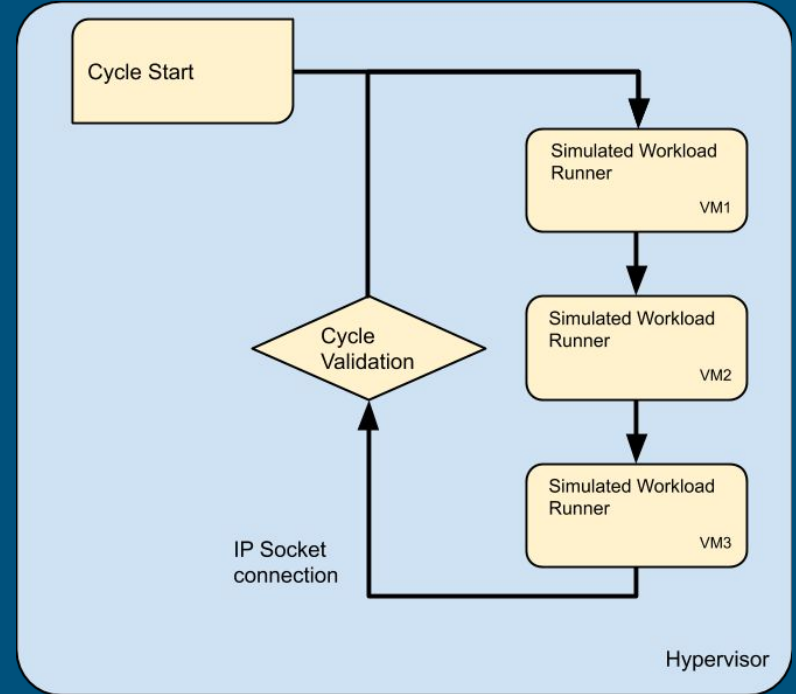
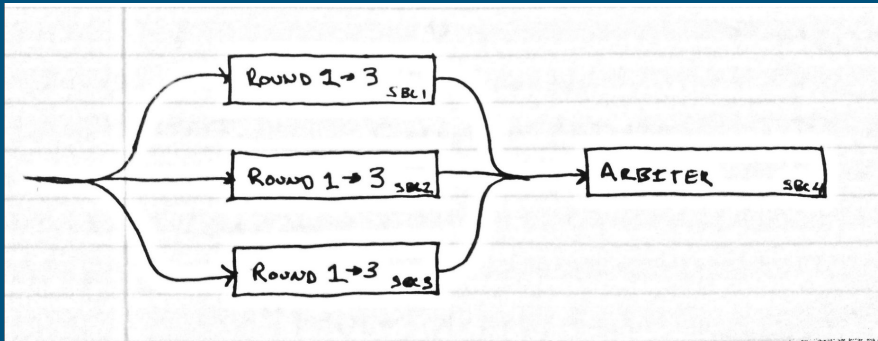

- Renewed global interest in space exploration
 - NASA Artemis
 - SpaceX Starship
 - Blue Origin Blue Moon



Test Plan

- Create a ecosystem of machines working together
- Inject errors into machines and observe error propagation
- Create different combinations of machines
- Evaluate the robustness of each combination, propagation of errors, etc

System Overview



System Details

- xinu00.cs.purdue.edu
- cs-console managing 185 Beaglebone Black ARM SBCs

```
# aneusted @ xinu17 in ~/ftd/SpaceWhale/Xinu-code-BeagleBoneBlack/compile on git:master x [22:46:19]
$ cs-status | head -n 5 && echo '...' && cs-status | tail -n 30
xinuserver:
  beagle101      cortex      user=          time=
  beagle102      cortex      user= hmonga   time= 484:41:13
  beagle103      cortex      user= dhar1     time= 151:52:29
  beagle104      cortex      user= cenrigh   time= 199:59:34
...
  beagle155      cortex      user=          time=
  beagle156      cortex      user=          time=
  beagle157      cortex      user=          time=
  beagle158      cortex      user=          time=
  beagle159      cortex      user=          time=
  beagle160      cortex      user= dkrolopp  time= 00:00:07
  beagle161      cortex      user= dkrolopp  time= 01:10:45
  beagle162      cortex      user= dkrolopp  time= 01:10:26
  beagle163      cortex      user=          time=
  beagle164      cortex      user=          time=
  beagle165      cortex      user=          time=
```

Machine Details

- BeagleBone Black SBCs running modified XINU operating system (ARM)
- Responsible for computing simulated workload and passing result to next SBC



Why Xinu?

- A small, embedded system much like those found on spacecraft
- Small memory footprint
- Easy fault instrumentation
- Familiarity

```
The text segment ranges from 0x81000000 to 0x810160a4 (Total size: 22569)
The data segment ranges from 0x81017000 to 0x81017868 (Total size: 538)
The BSS segment ranges from 0x81018000 to 0x810383e4 (Total size: 33017)
```

Simulated Workload Details

- Our simulated workload repeatedly sums numbers from an array in memory
 - A 4KB "workspace" of ints along with kernel memory serve as a target for faults
 - Virtually every fault in our workspace will cascade into a failure
 - Each round of summations is called an "iteration"
- Faults are instrumented programmatically from within the application
 - In our previous presentation, faults had a probability p of flipping a single bit after every iteration through our algorithm
 - With current design, a background process randomly schedules memory corruption events according to a uniform distribution

Faulty run results

```
finished test...  
Unhandlad exceptio.. Link Register: 0x||#  
  
panic: **** EXCEPTION ****
```

```
Unhandhad e8seption. \iji!Register:`0h!x  
  
panic: **** EXcEH\ION ****.
```

```
finished test...  
Unhandled exception. Link Register: 0x81007928  
  
panic: **** EXCEPTION ****
```

```
Received! Processing 100 iterations with input 1411548400...  
Processing complete. Final result: -1895994112.  
Wait time: 52611, compute time: 101807  
TEXT  flips: 65  
DATA  flips: 1  
BSS   flips: 85  
OTHER flips: 1  
finished test...
```

Xinu for bbb -- version #44 (aneusted) Thu Dec 12 23:12:16 EST 2019

Ethernet Link is Up. Speed is 100Mbps

Link is Full Duplex

MAC Address is: 6C:EC:EB:AC:58:EA

519731552 bytes of free memory. Free list:

[0x810526A0 to 0x9FFF9FFF]

89900 bytes of Xinu code.

[0x81000000 to 0x81015F2B]

136132 bytes of data.

[0x81016000 to 0x810373C3]

Obtained IP address 128.10.137.160 (0x800a89a0)

The text segment ranges from 0x81000000 to 0x81015f2c (Total size: 22475)

The data segment ranges from 0x81016000 to 0x81016868 (Total size: 538)

The BSS segment ranges from 0x81017000 to 0x810373c4 (Total size: 33009)

starting test...

Waiting on input...

Received! -133840484... Processing 100 iterations with input 314...

Processing complete. Sending result: -133840484...

Result sent!

Wait time: 33, compute time: 25483

TEXT flips: 0

DATA flips: 0

BSS flips: 2

OTHER flips: 0

finished test...

(command-mode) _

Xinu for bbb -- version #45 (aneusted) Thu Dec 12 23:18:25 EST 2019

Ethernet Link is Up. Speed is 100Mbps

Link is Full Duplex

MAC Address is: 6C:EC:EB:BA:1A:B7

519731552 bytes of free memory. Free list:

[0x810526A0 to 0x9FFF9FFF]

89896 bytes of Xinu code.

[0x81000000 to 0x81015F27]

136132 bytes of data.

[0x81016000 to 0x810373C3]

Obtained IP address 128.10.137.161 (0x800a89a1)

The text segment ranges from 0x81000000 to 0x81015f28 (Total size: 22474)

The data segment ranges from 0x81016000 to 0x81016868 (Total size: 538)

The BSS segment ranges from 0x81017000 to 0x810373c4 (Total size: 33009)

starting test...

Waiting on input...

Received! -133840484... Processing 100 iterations with input -133840484...

Processing complete. Sending result: 1411548400...

Result sent!

Wait time: 26543, compute time: 25548

TEXT flips: 23

DATA flips: 0

BSS flips: 30

OTHER flips: 0

finished test...

Xinu for bbb -- version #46 (aneusted) Thu Dec 12 23:18:44 EST 2019

Ethernet Link is Up. Speed is 100Mbps

Link is Full Duplex

MAC Address is: 6C:EC:EB:AB:07:8B

519731552 bytes of free memory. Free list:

[0x810526A0 to 0x9FFF9FFF]

89896 bytes of Xinu code.

[0x81000000 to 0x81015F27]

136132 bytes of data.

[0x81016000 to 0x810373C3]

Obtained IP address 128.10.137.162 (0x800a89a2)

The text segment ranges from 0x81000000 to 0x81015f28 (Total size: 22474)

The data segment ranges from 0x81016000 to 0x81016868 (Total size: 538)

The BSS segment ranges from 0x81017000 to 0x810373c4 (Total size: 33009)

starting test...

Waiting on input...

Received! Processing 100 iterations with input 1411548400...

Processing complete. Final result: -1895994112.

Wait time: 52611, compute time: 101807

TEXT flips: 65

DATA flips: 1

BSS flips: 85

OTHER flips: 1

finished test...


Data Collected

- Instrumented which node failures occurred, general cause, and errors injected in that node


O	P	Q	R	S	T	U	V	W	X	Y	Z
# TEXT bit flips (2)	# DATA bit flips (2)	# BSS bit flips (2)	# OTHER bit flips (2)	Total (2)	# TEXT bit flips (3)	# DATA bit flips (3)	# BSS bit flips (3)	# OTHER bit flips (3)	Total (3)	Outcome	Notes
0	0	0	0	0	0	0	0	0	0	Success	
0	0	2	0	2	-	-	-	-	0	Failure	Crash in node 2
17	1	35	0	53	-	-	-	-	0	Failure	Crash in node 3
21	0	30	1	52	48	1	55	1	105	Failure	Sent correct message 2 -> 3 but arrived differe
20	0	30	0	50	27	0	49	0	76	Success	
20	0	32	1	53	27	1	53	0	81	Success	
24	2	31	2	59	27	2	43	2	74	Success	
18	0	30	1	49	28	1	43	2	74	Failure	Incorrect computation happened in node 2
27	0	28	1	56	28	1	51	0	80	Success	
23	0	30	0	53	65	1	85	1	152	Failure	Node 4 took 4x as long as expected
15	0	33	2	50	31	1	53	0	85	Failure	Incorrect computation happened in node 3
22	0	26	2	50	32	0	42	3	77	Success	
20	0	20	1	41	26	0	42	1	69	Failure	Incorrect computation happened in node 2
25	0	32	0	57	-	-	-	-	0	Failure	Crash in node 3
21	0	35	0	56	30	-	48	2	80	Failure	Incorrect computation happened in node 2
0	0	0	0	0	-	-	-	-	0	Failure	Crash in node 2 during message sending
18.200	0.200	26.267	0.733	45.400	33.545	0.800	51.273	1.091	86.709		
40.09%	0.44%	57.86%	1.62%	100.00%	38.69%	0.92%	59.13%	1.26%	100.00%		

Analysis - Linear configuration

Outcomes observed

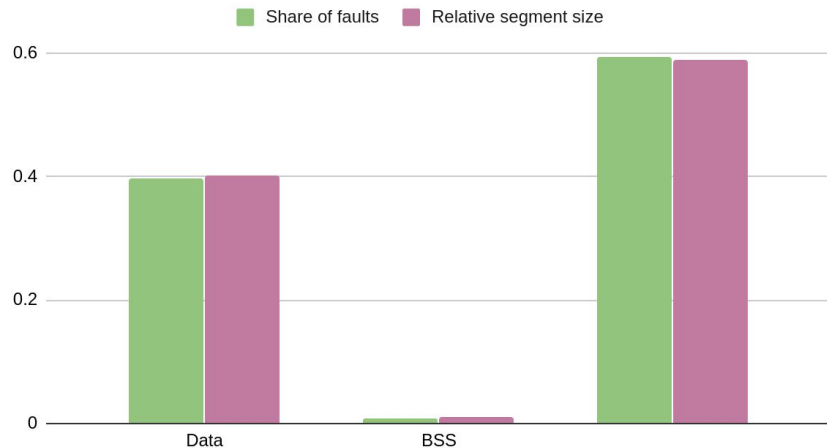


Average # of memory faults per node




Analysis - Linear configuration

Relative code segment size vs. share of faults




Which nodes caused failures?




Analysis - 2-of-3 Voting scheme

Average # of memory faults per node




Successful Computation Completions Per Node




Analysis - 2-of-3 Voting scheme

Node Outcomes observed



Number of Successful Nodes



Analysis - 2-of-3 Voting scheme

- The voting scheme actually performs significantly worse than the linear configuration
- At least 2 voters must perform correctly. The arbiter must also perform correctly
- As Prof. Bagchi put it, you're actually weaker than your weakest link

$$\binom{3}{2} \left(\frac{1}{3}\right)^2 \cdot \left(1 - \frac{1}{3}\right) \cdot \left(\frac{1}{3}\right) = 0.074$$

Issues

- Some data suspiciously repeats itself
 - Maybe RAM isn't being cleared as quickly as we expect it to be
- Lack of automation lead to limited data

Future Work

- Further analysis of collected data
- Additional testbed configurations
- Additional rates of fault injection
- Improvements in scalability and automation

Conclusion

- Even using a OS with a miniscule attack surface, operation is surprisingly robust
- Use caution before blindly adding “fault tolerate schemes” to failure prone systems